

# kubectl Cheat Sheet

This page contains a list of commonly used `kubectl` commands and flags.

## Kubectl autocomplete

### BASH

```
source <(kubectl completion bash) # set up autocomplete in bash into the current shell
echo "source <(kubectl completion bash)" >> ~/.bashrc # add autocomplete permanently to the profile
```

You can also use a shorthand alias for `kubectl` that also works with completion:

```
alias k=kubectl
complete -o default -F __start_kubectl k
```

### ZSH

```
source <(kubectl completion zsh) # set up autocomplete in zsh into the current shell
echo '[[ $commands[kubectl] ]] && source <(kubectl completion zsh)' >> ~/.zshrc # add autocomplete to the profile
```

## A note on `--all-namespaces`

Appending `--all-namespaces` happens frequently enough where you should be aware of the shorthand for `--all-namespaces` :

```
kubectl -A
```

## Kubectl context and configuration

Set which Kubernetes cluster `kubectl` communicates with and modifies configuration information. See [Authenticating Across Clusters with kubeconfig](#) documentation for detailed config file information.

```
kubectl config view # Show Merged kubeconfig settings.

# use multiple kubeconfig files at the same time and view merged config
KUBECONFIG=~/.kube/config:~/.kube/kubconfig2

kubectl config view

# get the password for the e2e user
kubectl config view -o jsonpath='{.users[?(@.name == "e2e")].user.password}'

kubectl config view -o jsonpath='{.users[].name}' # display the first user
```

```

kubectl config view -o jsonpath='{.users[*].name}'      # get a list of users
kubectl config get-contexts                            # display list of contexts
kubectl config current-context                        # display the current-context
kubectl config use-context my-cluster-name            # set the default context to my-cluster-name

kubectl config set-cluster my-cluster-name            # set a cluster entry in the kubeconfig file

# configure the URL to a proxy server to use for requests made by this client in the cluster
kubectl config set-cluster my-cluster-name --proxy-url=my-proxy-url

# add a new user to your kubeconfig that supports basic auth
kubectl config set-credentials kubeuser/foo.kubernetes.com --username=kubeuser --password=kubeuser

# permanently save the namespace for all subsequent kubectl commands in that context
kubectl config set-context --current --namespace=ggckad-s2

# set a context utilizing a specific username and namespace.
kubectl config set-context gce --user=cluster-admin --namespace=foo \
  && kubectl config use-context gce

kubectl config unset users.foo                      # delete user foo

# short alias to set/show context/namespace (only works for bash and bash-compatible shells)
alias kx='f() { [ "$1" ] && kubectl config use-context $1 || kubectl config current-context; }'
alias kn='f() { [ "$1" ] && kubectl config set-context --current --namespace $1 || kubectl config current-context; }'

```

## Kubectl apply

`apply` manages applications through files defining Kubernetes resources. It creates and updates resources in a cluster through running `kubectl apply`. This is the recommended way of managing Kubernetes applications on production. See [Kubectl Book](#).

## Creating objects

Kubernetes manifests can be defined in YAML or JSON. The file extension `.yaml`, `.yml`, and `.json` can be used.

```

kubectl apply -f ./my-manifest.yaml                  # create resource(s)
kubectl apply -f ./my1.yaml -f ./my2.yaml           # create from multiple files
kubectl apply -f ./dir                              # create resource(s) in all manifests in dir
kubectl apply -f https://git.io/vPieo                # create resource(s) from url
kubectl create deployment nginx --image=nginx        # start a single instance of nginx

# create a Job which prints "Hello World"
kubectl create job hello --image=busybox:1.28 -- echo "Hello World"

# create a CronJob that prints "Hello World" every minute
kubectl create cronjob hello --image=busybox:1.28 --schedule="*/1 * * * *" -- echo "Hello World"

kubectl explain pods                                # get the documentation for pod manifests

# Create multiple YAML objects from stdin
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox:1.28
    args:
    - sleep
    - "1000000"

```

```

---
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep-less
spec:
  containers:
  - name: busybox
    image: busybox:1.28
    args:
    - sleep
    - "1000"
EOF

# Create a secret with several keys
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: $(echo -n "s33msi4" | base64 -w0)
  username: $(echo -n "jane" | base64 -w0)
EOF

```

## Viewing and finding resources

```

# Get commands with basic output
kubectl get services                                # List all services in the namespace
kubectl get pods --all-namespaces                   # List all pods in all namespaces
kubectl get pods -o wide                           # List all pods in the current namespace
kubectl get deployment my-dep                      # List a particular deployment
kubectl get pods                                    # List all pods in the namespace
kubectl get pod my-pod -o yaml                     # Get a pod's YAML

# Describe commands with verbose output
kubectl describe nodes my-node
kubectl describe pods my-pod

# List Services Sorted by Name
kubectl get services --sort-by=.metadata.name

# List pods Sorted by Restart Count
kubectl get pods --sort-by='.status.containerStatuses[0].restartCount'

# List PersistentVolumes sorted by capacity
kubectl get pv --sort-by=.spec.capacity.storage

# Get the version label of all pods with label app=cassandra
kubectl get pods --selector=app=cassandra -o \
  jsonpath='{.items[*].metadata.labels.version}'

# Retrieve the value of a key with dots, e.g. 'ca.crt'
kubectl get configmap myconfig \
  -o jsonpath='{.data.ca\.crt}'

# Retrieve a base64 encoded value with dashes instead of underscores.
kubectl get secret my-secret --template='{{index .data "key-name-with-dashes"}}'

# Get all worker nodes (use a selector to exclude results that have a label
# named 'node-role.kubernetes.io/control-plane')
kubectl get node --selector='!node-role.kubernetes.io/control-plane'

# Get all running pods in the namespace
kubectl get pods --field-selector=status.phase=Running

```

```

# Get ExternalIPs of all nodes
kubectl get nodes -o jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP")].address}'

# List Names of Pods that belong to Particular RC
# "jq" command useful for transformations that are too complex for jsonpath, it can be used like this
sel=${$(kubectl get rc my-rc --output=json | jq -j '.spec.selector | to_entries | .[0].value')}
echo $(kubectl get pods --selector=$sel --output=jsonpath='{.items..metadata.name}')

# Show labels for all pods (or any other Kubernetes object that supports labelling)
kubectl get pods --show-labels

# Check which nodes are ready
JSONPATH='{range .items[*]}{@.metadata.name}:{range @.status.conditions[*]}{@.type}:{@.status}'
&& kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"

# Output decoded secrets without external tools
kubectl get secret my-secret -o go-template='{{range $k,$v := .data}}{{"### "}}{{ $k }}: {{ $v | base64decode }}\n{{end}}'

# List all Secrets currently in use by a pod
kubectl get pods -o json | jq '.items[].spec.containers[].env[]?.valueFrom.secretName'

# List all containerIDs of initContainer of all pods
# Helpful when cleaning up stopped containers, while avoiding removal of initContainers
kubectl get pods --all-namespaces -o jsonpath='{range .items[*]}{@.metadata.namespace}:{@.status.initContainerID}\n{{end}}'

# List Events sorted by timestamp
kubectl get events --sort-by=.metadata.creationTimestamp

# Compares the current state of the cluster against the state that the cluster would have been in if the given manifest was the only one applied
kubectl diff -f ./my-manifest.yaml

# Produce a period-delimited tree of all keys returned for nodes
# Helpful when locating a key within a complex nested JSON structure
kubectl get nodes -o json | jq -c 'paths|join(".")'

# Produce a period-delimited tree of all keys returned for pods, etc
kubectl get pods -o json | jq -c 'paths|join(".")'

# Produce ENV for all pods, assuming you have a default container for the pods, and you want to run a command across all pods, not just `env`
for pod in $(kubectl get po --output=jsonpath='{.items..metadata.name}'); do echo $pod: $(kubectl exec $pod -- env) \n; done

# Get a deployment's status subresource
kubectl get deployment nginx-deployment --subresource=status

```

## Updating resources

```

kubectl set image deployment/frontend www=image:v2 # Rolling update
kubectl rollout history deployment/frontend # Check the history
kubectl rollout undo deployment/frontend # Rollback to the previous version
kubectl rollout undo deployment/frontend --to-revision=2 # Rollback to a specific revision
kubectl rollout status -w deployment/frontend # Watch rolling
kubectl rollout restart deployment/frontend # Rolling restart

cat pod.json | kubectl replace -f - # Replace a pod

# Force replace, delete and then re-create the resource. Will cause a service outage
kubectl replace --force -f ./pod.json

# Create a service for a replicated nginx, which serves on port 80 and connects to the backend pods on port 8080
kubectl expose rc nginx --port=80 --target-port=8080

# Update a single-container pod's image version (tag) to v4
kubectl get pod mypod -o yaml | sed 's/\(image: myimage\):.*$/\1:v4/' | kubectl replace -f -

```

```
kubectl label pods my-pod new-label=awesome           # Add a Label
kubectl label pods my-pod new-label-                  # Remove a label
kubectl annotate pods my-pod icon-url=http://goo.gl/XXBTWq  # Add an annotation
kubectl autoscale deployment foo --min=2 --max=10       # Auto scale a deployment
```

## Patching resources

```
# Partially update a node
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}'

# Update a container's image; spec.containers[*].name is required because it's a
kubectl patch pod valid-pod -p '{"spec":{"containers":[{"name":"kubernetes-serve-

# Update a container's image using a json patch with positional arrays
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path": "/spec/c

# Disable a deployment livenessProbe using a json patch with positional arrays
kubectl patch deployment valid-deployment --type json -p='[{"op": "remove", "p

# Add a new element to a positional array
kubectl patch sa default --type='json' -p='[{"op": "add", "path": "/secrets/1", "

# Update a deployment's replica count by patching its scale subresource
kubectl patch deployment nginx-deployment --subresource='scale' --type='merge' -p
```

## Editing resources

Edit any API resource in your preferred editor.

```
kubectl edit svc/docker-registry           # Edit the service named docker-registry
KUBE_EDITOR="nano" kubectl edit svc/docker-registry  # Use an alternative editor
```

## Scaling resources

```
kubectl scale --replicas=3 rs/foo           # Scale a replication controller
kubectl scale --replicas=3 -f foo.yaml       # Scale a resource from a file
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql  # If the deployment
kubectl scale --replicas=5 rc/foo rc/bar rc/baz  # Scale multiple resources
```

## Deleting resources

```
kubectl delete -f ./pod.json                # Delete a pod from a manifest file
kubectl delete pod unwanted --now           # Delete a pod immediately
kubectl delete pod,service baz foo         # Delete pods and services named baz
kubectl delete pods,services -l name=myLabel  # Delete pods and services with label
kubectl -n my-ns delete pod,svc --all       # Delete all pods and services in namespace
# Delete all pods matching the awk pattern1 or pattern2
kubectl get pods -n mynamespace --no-headers=true | awk '/pattern1|pattern2/{pri
```

# Interacting with running Pods

```
kubectl logs my-pod # dump pod logs (stdout)
kubectl logs -l name=myLabel # dump pod logs, with label n
kubectl logs my-pod --previous # dump pod logs (stdout) for
kubectl logs my-pod -c my-container # dump pod container logs (st
kubectl logs -l name=myLabel -c my-container # dump pod logs, with label n
kubectl logs my-pod -c my-container --previous # dump pod container logs (st
kubectl logs -f my-pod # stream pod logs (stdout)
kubectl logs -f my-pod -c my-container # stream pod container logs (
kubectl logs -f -l name=myLabel --all-containers # stream all pods logs with l
kubectl run -i --tty busybox --image=busybox:1.28 -- sh # Run pod as interactive
kubectl run nginx --image=nginx -n mynamespace # Start a single instance of
kubectl run nginx --image=nginx --dry-run=client -o yaml > pod.yaml # Generate spec for running p
kubectl attach my-pod -i # Attach to Running Container
kubectl port-forward my-pod 5000:6000 # Listen on port 5000 on the
kubectl exec my-pod -- ls / # Run command in existing pod
kubectl exec --stdin --tty my-pod -- /bin/sh # Interactive shell access to
kubectl exec my-pod -c my-container -- ls / # Run command in existing pod
kubectl top pod POD_NAME --containers # Show metrics for a given po
kubectl top pod POD_NAME --sort-by=cpu # Show metrics for a given po
```

# Copying files and directories to and from containers

```
kubectl cp /tmp/foo_dir my-pod:/tmp/bar_dir # Copy /tmp/foo_dir local
kubectl cp /tmp/foo my-pod:/tmp/bar -c my-container # Copy /tmp/foo local file
kubectl cp /tmp/foo my-namespace/my-pod:/tmp/bar # Copy /tmp/foo local file
kubectl cp my-namespace/my-pod:/tmp/foo /tmp/bar # Copy /tmp/foo from a rem
```

**Note:** `kubectl cp` requires that the 'tar' binary is present in your container image. If 'tar' is not present, `kubectl cp` will fail. For advanced use cases, such as symlinks, wildcard expansion or file mode preservation consider using `kubectl exec`.

```
tar cf - /tmp/foo | kubectl exec -i -n my-namespace my-pod -- tar xf - -C /tmp/ba
kubectl exec -n my-namespace my-pod -- tar cf - /tmp/foo | tar xf - -C /tmp/bar
```

# Interacting with Deployments and Services

```
kubectl logs deploy/my-deployment # dump Pod logs for a D
kubectl logs deploy/my-deployment -c my-container # dump Pod logs for a D

kubectl port-forward svc/my-service 5000 # listen on local port
kubectl port-forward svc/my-service 5000:my-service-port # listen on local port

kubectl port-forward deploy/my-deployment 5000:6000 # listen on local port
kubectl exec deploy/my-deployment -- ls # run command in first
```



# Interacting with Nodes and cluster

```
kubectl cordon my-node           # Mark my-node as unschedulable
kubectl drain my-node            # Drain my-node
kubectl uncordon my-node         # Mark my-node as schedulable
kubectl top node my-node         # Show metrics for a node
kubectl cluster-info             # Display cluster information
kubectl cluster-info dump        # Dump cluster information
kubectl cluster-info dump --output-directory=/path/to/cluster-state # Dump cluster information to a directory

# View existing taints on which exist on current nodes.
kubectl get nodes -o='custom-columns=NodeName:.metadata.name,TaintKey:.spec.taints[*].key,TaintValue:.spec.taints[*].value'

# If a taint with that key and effect already exists, its value is replaced as specified
kubectl taint nodes foo dedicated=special-user:NoSchedule
```

## Resource types

List all supported resource types along with their shortnames, [API group](#), whether they are [namespaced](#), and [Kind](#):

```
kubectl api-resources
```

Other operations for exploring API resources:

```
kubectl api-resources --namespaced=true      # All namespaced resources
kubectl api-resources --namespaced=false     # All non-namespaced resources
kubectl api-resources -o name                 # All resources with simple output (name)
kubectl api-resources -o wide                 # All resources with expanded (aka "wide") output
kubectl api-resources --verbs=list,get        # All resources that support the "list" and "get" verbs
kubectl api-resources --api-group=extensions # All resources in the "extensions" API group
```

## Formatting output

To output details to your terminal window in a specific format, add the `-o` (or `--output`) flag to a supported `kubectl` command.

Output format	Description
<code>-o=custom-columns=&lt;spec&gt;</code>	Print a table using a comma separated list of custom columns
<code>-o=custom-columns-file=&lt;filename&gt;</code>	Print a table using the custom columns template in the <code>&lt;filename&gt;</code> file
<code>-o=json</code>	Output a JSON formatted API object
<code>-o=jsonpath=&lt;template&gt;</code>	Print the fields defined in a <a href="#">jsonpath</a> expression
<code>-o=jsonpath-file=&lt;filename&gt;</code>	Print the fields defined by the <a href="#">jsonpath</a> expression in the <code>&lt;filename&gt;</code> file
<code>-o=name</code>	Print only the resource name and nothing else

Output format	Description
-o=wide	Output in the plain-text format with any additional information, and for pods, the node name is included
-o=yaml	Output a YAML formatted API object

Examples using `-o=custom-columns` :

```
# All images running in a cluster
kubectl get pods -A -o=custom-columns='DATA:spec.containers[*].image'

# All images running in namespace: default, grouped by Pod
kubectl get pods --namespace default --output=custom-columns="NAME:.metadata.name,DATA:spec.containers[*].image"

# All images excluding "registry.k8s.io/coredns:1.6.2"
kubectl get pods -A -o=custom-columns='DATA:spec.containers[?(@.image!="registry.k8s.io/coredns:1.6.2")].image'

# All fields under metadata regardless of name
kubectl get pods -A -o=custom-columns='DATA:metadata.*'
```

More examples in the kubectl [reference documentation](#).

## Kubectl output verbosity and debugging

Kubectl verbosity is controlled with the `-v` or `--v` flags followed by an integer representing the log level. General Kubernetes logging conventions and the associated log levels are described [here](#).

Verbosity	Description
- -v=0	Generally useful for this to <i>always</i> be visible to a cluster operator.
- -v=1	A reasonable default log level if you don't want verbosity.
- -v=2	Useful steady state information about the service and important log messages that may correlate to significant changes in the system. This is the recommended default log level for most systems.
- -v=3	Extended information about changes.
- -v=4	Debug level verbosity.
- -v=5	Trace level verbosity.
- -v=6	Display requested resources.
- -v=7	Display HTTP request headers.
- -v=8	Display HTTP request contents.

- -v=9      Display HTTP request contents without truncation of contents.

## What's next

- Read the [kubectl overview](#) and learn about [JsonPath](#).



- See [kubectl](#) options.
- Also read [kubectl Usage Conventions](#) to understand how to use kubectl in reusable scripts.
- See more community [kubectl cheatsheets](#).

## Feedback

Was this page helpful?

Yes

No

---

Last modified November 22, 2022 at 9:21 PM PST: [Fix typos in cheatsheet.md \(3e20970eb3\)](#)